

Introduction to Information Retrieval

Lecture 6: Scoring, Term Weighting and the Vector Space Model

Contents

- Parametric and zone indexes
- Ranked retrieval
- Term Weighting schemes
- Vector space scoring

Parametric and zone indexes

- Each document has, in addition to text, some “meta-data” in fields e.g.,
 - Language = French
 - Format = pdf
 - Subject = Physics etc.
 - Date = Feb 2000
- A parametric search interface allows the user to combine a full-text query with selections on these field values e.g.,
 - language, date range, etc.

Fields

Values

User view

The screenshot shows a 'Bibliographic Search' form with the following fields and options:

- Search criteria:** A dropdown menu.
- Author:** A text input field with the example 'Example: Wilson, J or Einstein-Milner'.
- Title:** A text input field with the instruction 'Also a part of the title possible'.
- Date of publication:** A text input field with the example 'Example: 1997 or 1997- or 1997- limits the search to the documents appeared in, before and after 1997 respectively'.
- Language:** A dropdown menu with the instruction 'Language the document was written in' and the selected value 'English (3)'.
- Project:** A dropdown menu with the value 'ANY'.
- Type:** A dropdown menu with the value 'ANY'.
- Subject group:** A dropdown menu with the value 'ANY'.
- Sorted by:** A dropdown menu with the value 'Date of publication'.
- Buttons:** 'Start bibliographic search' and 'Find document via ID'.

► Figure 6.1 Parametric search. In this example we have a collection with fields allowing us to select publications by zones such as Author and fields such as Language.

4

Zones

- A zone is an identified region within a doc
 - E.g., Title, Abstract, Bibliography
- Contents of a zone are free text
 - Not a “finite” vocabulary
- Indexes for each zone - allow queries like
 - find documents with merchant in the title and william in the author list and the phrase gentle rain in the body

Zone indexes – simple view

Term	IN doc id	Term	IN doc id	Term	IN doc id
merchant	1	william	1	gentle	1
rain	1	in	1	the	1
title	1	author	1	body	1
abstract	1	body	1	etc.	1
etc.	1				

Title Author Body

Introduction to Information Retrieval

Zone indexes – common view

► Figure 6.2 Basic zone index ; zones are encoded as extensions of dictionary entries.

► Figure 6.3 Zone index in which the zone is encoded in the postings rather than the dictionary.

7

Introduction to Information Retrieval

Weighted zone scoring

- Given a Boolean query q and a document d , *weighted zone scoring* assigns to the pair (q, d) a score in the interval $[0, 1]$,
 - by computing a linear combination of zone scores
 - where each zone of the document contributes a Boolean value.
- Specifically,
 - let there are L zones. Let $g_1, \dots, g_L \in [0, 1]$ such that $\sum_{i=1}^L g_i = 1$
 - let s be the Boolean score denoting a match (or absence thereof) between q and the i th zone.
 - Then, the weighted zone score is defined to be $\sum_{i=1}^L g_i \cdot s_i$

8

Introduction to Information Retrieval

Example: Weighted zone scoring

- Query $Q = \text{shakespeare}$
- consider a collection in which each document has three zones: **author, title and body**
- Suppose we set $g_1 = 0.2$, $g_2 = 0.3$ and $g_3 = 0.5$ where g_1 , g_2 and g_3 represents the **author, title and body zone weights**.
- If the term shakespeare appear in the **title** and **body** zones but not the **author** zone of a document, the score of this document would be
- 0.8.

9

Introduction to Information Retrieval

Ranked retrieval models

- Rather than a set of documents satisfying a query expression, in ranked retrieval, the system returns an ordering over the (top) documents in the collection for a query
- Free text queries: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language

10

Introduction to Information Retrieval

Sec. 6.2

Recall (Lecture 1): Binary term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

Introduction to Information Retrieval

Sec. 6.2

Term-document count matrices

- Consider the number of occurrences of a term in a document:
 - Each document is a count vector: a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Introduction to Information Retrieval

Bag of words model

- Vector representation doesn't consider the ordering of words in a document
- John is quicker than Mary* and *Mary is quicker than John* have the same vectors
- This is called the bag of words model.
- In a sense, this is a step back: The positional index was able to distinguish these two documents.
- We will look at "recovering" positional information later in this course.
- For now: bag of words model

Introduction to Information Retrieval

Term frequency tf

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- We want to use tf when computing query-document match scores. But how?
- Raw term frequency is not what we want:
 - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
 - But not 10 times more relevant.
- Relevance does not increase proportionally with term frequency.

NB: frequency = count in IR

Introduction to Information Retrieval Sec. 6.2

Log-frequency weighting

- The log frequency weight of term t in d is

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d}, & \text{if } tf_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$
- $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$, etc.
- Score for a document-query pair: sum over terms t in both q and d .
- score = $\sum_{t \in q \cap d} (1 + \log tf_{t,d})$
- The score is 0 if none of the query terms is present in the document.

Introduction to Information Retrieval Sec. 6.2.1

Document frequency

- Rare terms are more informative than frequent terms
 - Recall stop words
- We want a high weight for rare terms.

Introduction to Information Retrieval Sec. 6.2.1

idf weight

- df_t is the document frequency of t : the number of documents that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the idf (inverse document frequency) of t by

$$idf_t = \log_{10} (N/df_t)$$
 - We use $\log (N/df_t)$ instead of N/df_t to "dampen" the effect of idf.

Will turn out the base of the log is immaterial.

Introduction to Information Retrieval Sec. 6.2.1

idf example, suppose $N = 806791$

term	df_t	idf_t
car	18,165	1.65
auto	6723	2.08
insurance	19,241	1.62
best	25,235	1.5

$$idf_t = \log_{10} (N/df_t)$$

There is one idf value for each term t in a collection.

Introduction to Information Retrieval Sec. 6.2.2

tf-idf weighting

- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$\text{tf-idf}_{t,d} = \log(1 + \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$
- Best known weighting scheme in information retrieval
 - Note: the "-" in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf
- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

Introduction to Information Retrieval Sec. 6.3

Binary → count → tf-idf weight matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worse	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf

Introduction to Information Retrieval Sec. 6.2.2

Score for a document given a query

$$\text{Score}(q, d) = \sum_{t \in q \cap d} \text{tf.idf}_{t,d}$$

- There are many variants
 - How "tf" is computed (with/without logs)
 - Whether the terms in the query are also weighted
 - ...

21

Introduction to Information Retrieval Sec. 6.3

The Vector Space Model for Scoring

- The set of documents in a collection may be viewed as a set of vectors in a vector space.
 - Terms are axes of the space
 - Documents are points or vectors in this space
- Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- These are very sparse vectors - most entries are zero.

Introduction to Information Retrieval

Vector Similarity

- How do we quantify the similarity between two documents in this vector space?
- A first attempt : the magnitude of the vector difference between two document vectors.
 - Drawback:** two documents with very similar content can have a significant vector difference simply because one is much longer than the other.
- Solution** to compensate for the effect of document length is to compute the *cosine similarity*.

23

Introduction to Information Retrieval

Cosine similarity illustrated

24

Introduction to Information Retrieval Sec. 6.3

From angles to cosines

- The following two notions are equivalent.
 - Rank documents in decreasing order of the angle between query and document
 - Rank documents in increasing order of $\text{cosine}(\text{query}, \text{document})$
- Cosine is a monotonically decreasing function for the interval $[0^\circ, 180^\circ]$

Introduction to Information Retrieval Sec. 6.3

$\text{cosine}(\text{document}, \text{document})$

DOT PRODUCT

$$\text{sim}(d_1, d_2) = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| |\vec{V}(d_2)|} \quad (2)$$

EUCLIDEAN LENGTH

The dot product $\vec{x} \cdot \vec{y}$ of two vectors is defined as $\sum_{i=1}^M x_i y_i$
 The Euclidean length of d is defined to be $\sqrt{\sum_{i=1}^M \vec{V}_i^2(d)}$.

Introduction to Information Retrieval Sec. 6.3

Cosine for length-normalized vectors

The effect of the denominator of Equation (2) is thus to *length-normalize* the vectors $\vec{V}(d_1)$ and $\vec{V}(d_2)$ to unit vectors $\vec{v}(d_1) = \vec{V}(d_1)/|\vec{V}(d_1)|$ and $\vec{v}(d_2) = \vec{V}(d_2)/|\vec{V}(d_2)|$

Then we can rewrite the previous equation as:

$$\text{sim}(d_1, d_2) = \vec{v}(d_1) \cdot \vec{v}(d_2). \quad (3)$$

27

Introduction to Information Retrieval Sec. 6.4

Example: N= 1,000,000

Document: *car insurance auto insurance*
Query: *best car insurance*

Term	DF
auto	5000
best	50000
car	10000
insurance	1000

Introduction to Information Retrieval Sec. 6.4

Example: N= 1,000,000

Document: *car insurance auto insurance*
Query: *best car insurance*

Term	Query	Document	Pro d
auto			
best			
car			
insurance			

Introduction to Information Retrieval Sec. 6.4

Example: N= 1,000,000

Document: *car insurance auto insurance*
Query: *best car insurance*

Term	Query						Document				Pro d
	tf-raw	tf-wt	df	idf	wt	n'lize	tf-raw	tf-wt	wt	n'lize	
auto	0	0	5000	2.3	0	0	1	1	2.3	0.47	0
best	1	1	50000	1.3	1.3	0.34	0	0	0	0	0
car	1	1	10000	2.0	2.0	0.53	1	1	2.0	0.41	0.22
insurance	1	1	1000	3.0	3.0	0.79	2	1.3	3.9	0.80	0.63

Query length = $\sqrt{0^2 + 1.3^2 + 2^2 + 3^2} \approx 3.8$
 Doc length = $\sqrt{2.3^2 + 0^2 + 2^2 + 3.9^2} \approx 4.9$
 Score = $0+0+0.22+0.63 = 0.85$

Introduction to Information Retrieval Sec. 6.4

Example: $N = 1,000,000$

Document: *car insurance auto insurance*

Query: *best car insurance*

Term	Query					Document					Prod
	tf-raw	tf-wt	df	idf	wt	n'lize	tf-raw	tf-wt	wt	n'lize	
auto	0	0	5000	2.3	0	0	1	1	2.3	0.47	0
best	1	1	50000	1.3	1.3	0.34	0	0	0	0	0
car	1	1	10000	2.0	2.0	0.53	1	1	2.0	0.41	0.22
insurance	1	1	1000	3.0	3.0	0.79	2	1.3	3.9	0.80	0.63

Query length = $\sqrt{0^2 + 1.3^2 + 2^2 + 3^2} \approx 3.8$

Doc length = $\sqrt{2.3^2 + 0^2 + 2^2 + 3.9^2} \approx 4.9$

Score = $0 + 0 + 0.22 + 0.63 = 0.85$

Introduction to
Information Retrieval

Lecture 8: Evaluation IR systems

Introduction to Information Retrieval Sec. 8.6

Measures for a search engine

- How fast does it index
 - Number of documents/hour
 - (Average document size)
- How fast does it search
 - Latency as a function of index size
- Quality of results
 - Precision
 - Recall
 - F-measure
- Expressiveness of query language
 - Ability to express complex information needs

Efficiency (covers index and search speed)

Effectiveness (covers quality of results)

Usability (covers expressiveness)

33

Introduction to Information Retrieval Sec. 8.1

Evaluating an IR system

- Note: the **information need** is translated into a **query**
- Relevance is assessed relative to the **information need** *not* the **query**
- E.g., **Information need:** *I'm looking for information on whether travelling by train from Cairo to Assuit is more effective than flying.*
- Query:** *travelling by train from Cairo to Assuit effective*
- Evaluate whether the doc addresses the information need, not whether it has these words

34

Introduction to Information Retrieval Sec. 8.2

Standard relevance benchmarks

- TREC - National Institute of Standards and Technology (NIST) has run a large IR test bed for many years
- Reuters and other benchmark doc collections used
- "Retrieval tasks" specified
 - sometimes as queries
- Human experts mark, for each query and for each doc, **Relevant** or **Nonrelevant**
 - or at least for subset of docs that some system returned for that query

35

Introduction to Information Retrieval Sec. 8.3

Unranked retrieval evaluation: Precision and Recall

- Precision:** fraction of retrieved docs that are relevant = (relevant retrieved / retrieved)
- Recall:** fraction of relevant docs that are retrieved = (relevant retrieved / relevant)

	Relevant	Nonrelevant
Retrieved	tp	fp
Not Retrieved	fn	tn

- Precision $P = tp / (tp + fp)$
- Recall $R = tp / (tp + fn)$

36

Introduction to Information Retrieval Sec. 8.3

Should we instead use the accuracy measure for evaluation?

- Given a query, an engine classifies each doc as “Relevant” or “Nonrelevant”
- The **accuracy** of an engine: the fraction of these classifications that are correct
 - $(tp + tn) / (tp + fp + fn + tn)$
- Accuracy** is a commonly used evaluation measure in machine learning classification work

37

Introduction to Information Retrieval Sec. 8.3

Precision/Recall

- You can get **high recall** (but low precision) by **retrieving all docs for all queries!**
- Recall is a non-decreasing function of the number of docs retrieved
- In a good system, precision decreases as either the number of docs retrieved or recall increases
 - This is not a theorem, but a result with strong empirical confirmation

38

Introduction to Information Retrieval Sec. 8.3

A combined measure: F

- Combined measure that assesses precision/recall tradeoff is **F measure** (weighted harmonic mean):

$$F = \frac{1}{\alpha \frac{1}{P} + (1-\alpha) \frac{1}{R}}$$
- People usually use balanced F_1 measure
 - i.e., with $\alpha = \frac{1}{2}$
$$F_1 = \frac{2PR}{P + R}$$

39

Introduction to Information Retrieval Sec. 8.4

Evaluating ranked results

- Evaluation of ranked results:
 - The system can return any number of results
 - By taking various numbers of the top returned documents (levels of recall), the evaluator can produce a *precision-recall curve*

40

Introduction to Information Retrieval Sec. 8.4

Averaging over queries

- A precision-recall graph for one query isn’t a very sensible thing to look at
- You need to average performance over a whole bunch of queries.
- But there’s a technical issue:
 - Precision-recall calculations place some points on the graph

41

Introduction to Information Retrieval Sec. 8.4

Typical (good) 11 point precisions

42